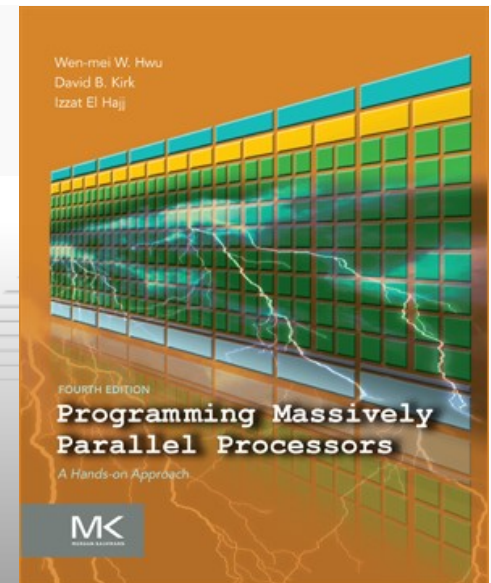
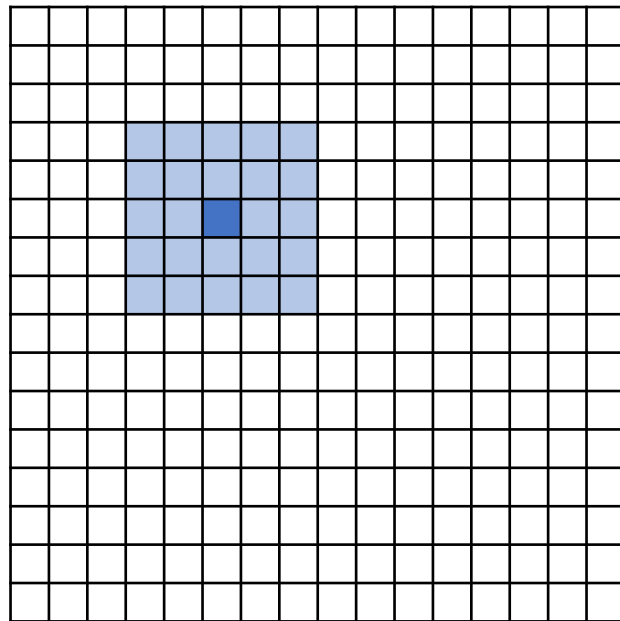


Programming Massively Parallel Processors

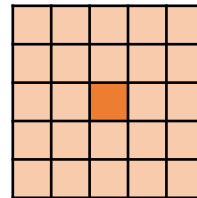
A Hands-on Approach

CHAPTER 7 > Convolution

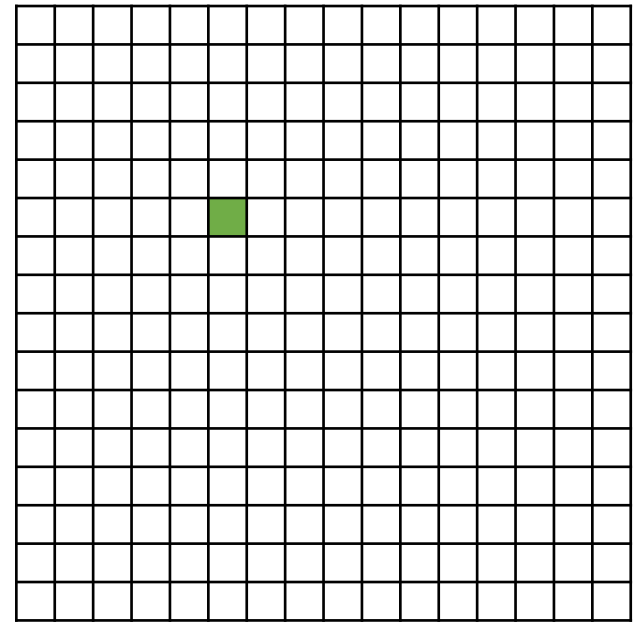




input



filter



output

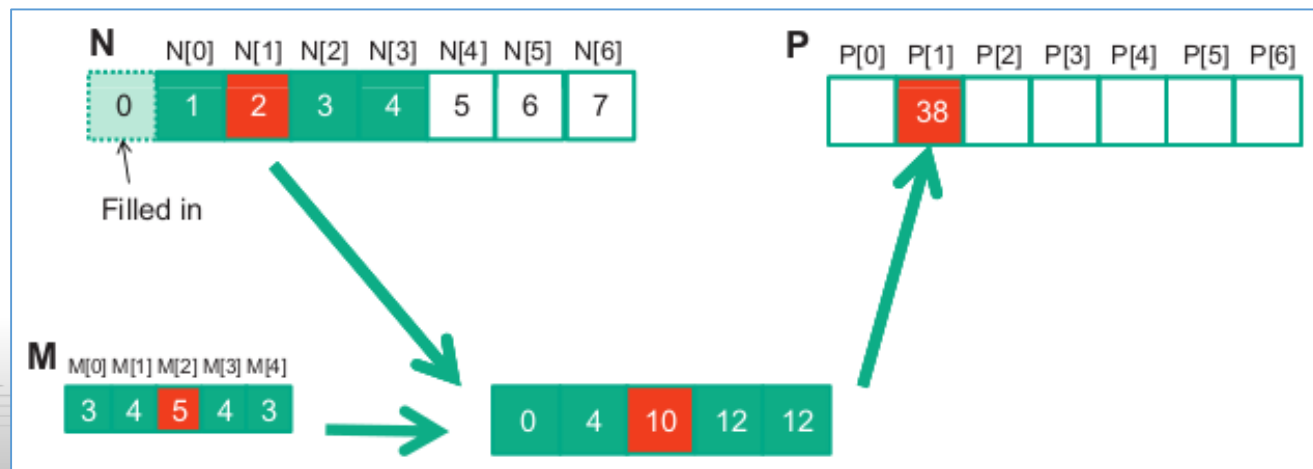
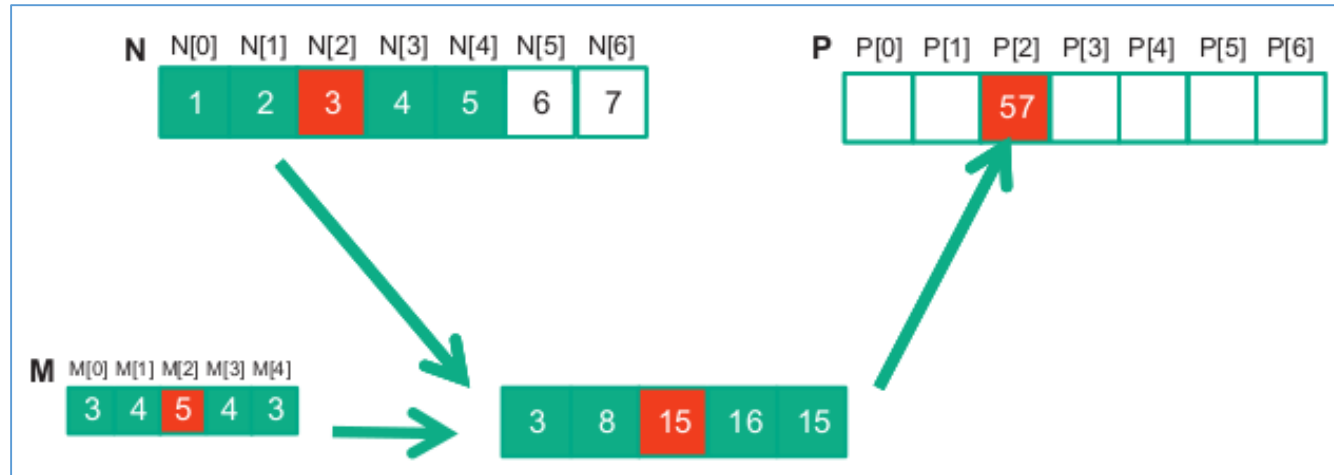
Every **output element** is a weighted sum of the neighboring **input elements**

Image blur seen before was a special case where all weights are the same

In general, weights are determined by a convolution **filter**

(commonly called convolution *kernel*, but we will use *filter* to avoid confusion with CUDA kernels)

- Convolution is an array operation where each output data element is a weighted sum of a collection of neighboring input elements.



N

1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

P

		321				


M

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1




1	4	9	8	5
4	9	16	15	12
9	16	25	24	21
8	15	24	21	16
5	12	21	16	5

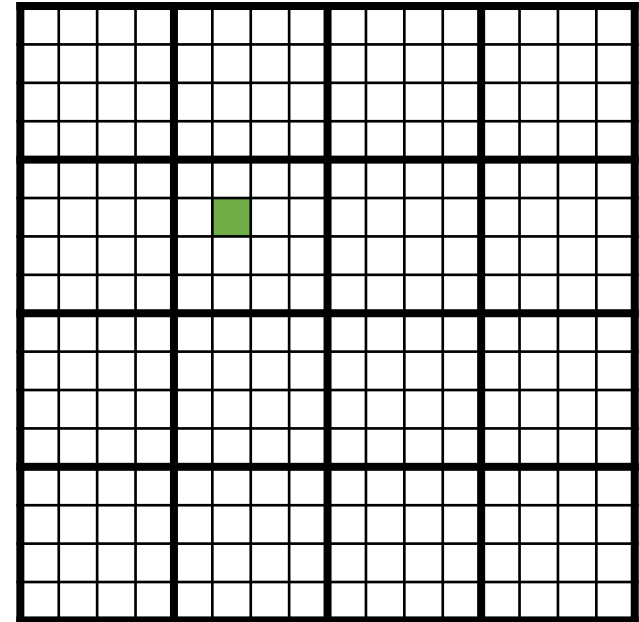
- Commonly used in signal processing, image processing, video processing, etc.
- Usually used to transform signals or pixels to more desirable values
 - e.g., Gaussian blur, sharpen, edge detection, etc.
 - Transformation depends on the weights in the filter
- Using 2D as an example, but can also be 1D or 3D



input



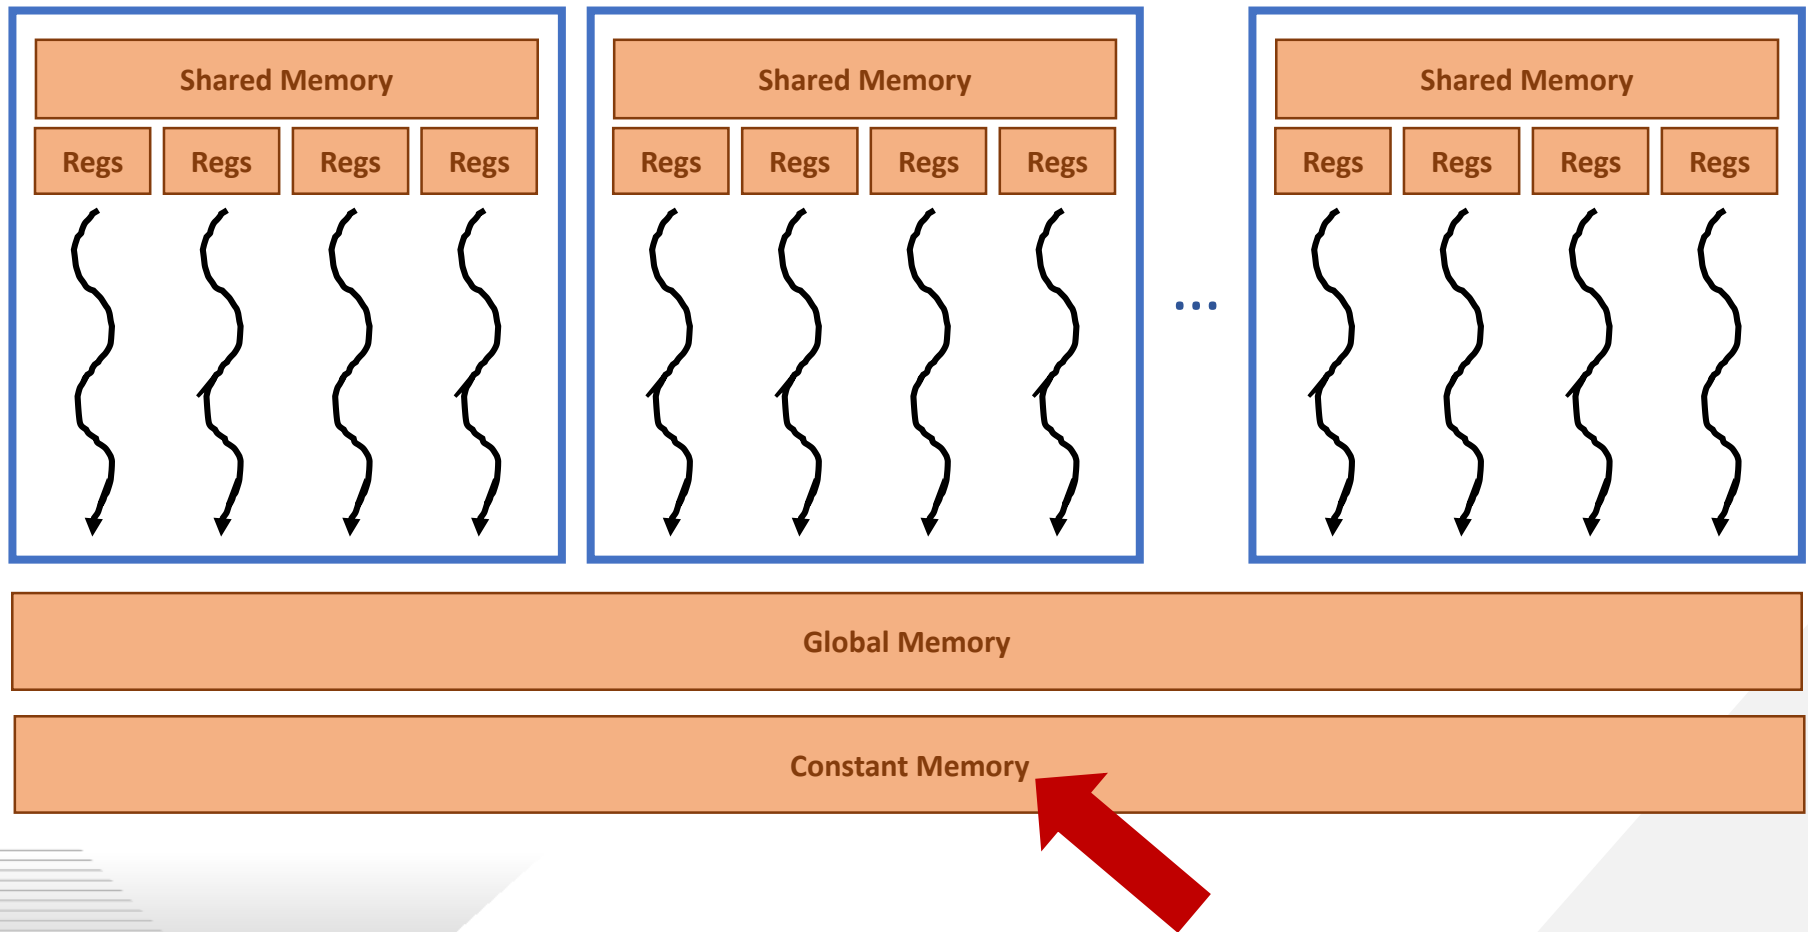
filter



output

Parallelization approach: Assign one thread to compute each **output element** by looping over **input elements** and **filter** weights

- Observations:
 - The filter is typically small
 - The filter is constant (weights do not change)
 - The filter is accessed by all threads in the grid
- Optimization: store the filter in **constant memory** for quicker access



- Declare constant memory array as global variable

```
__constant__ float filter_c[FILTER_DIM][FILTER_DIM];
```

- Must initialize constant memory from the host:

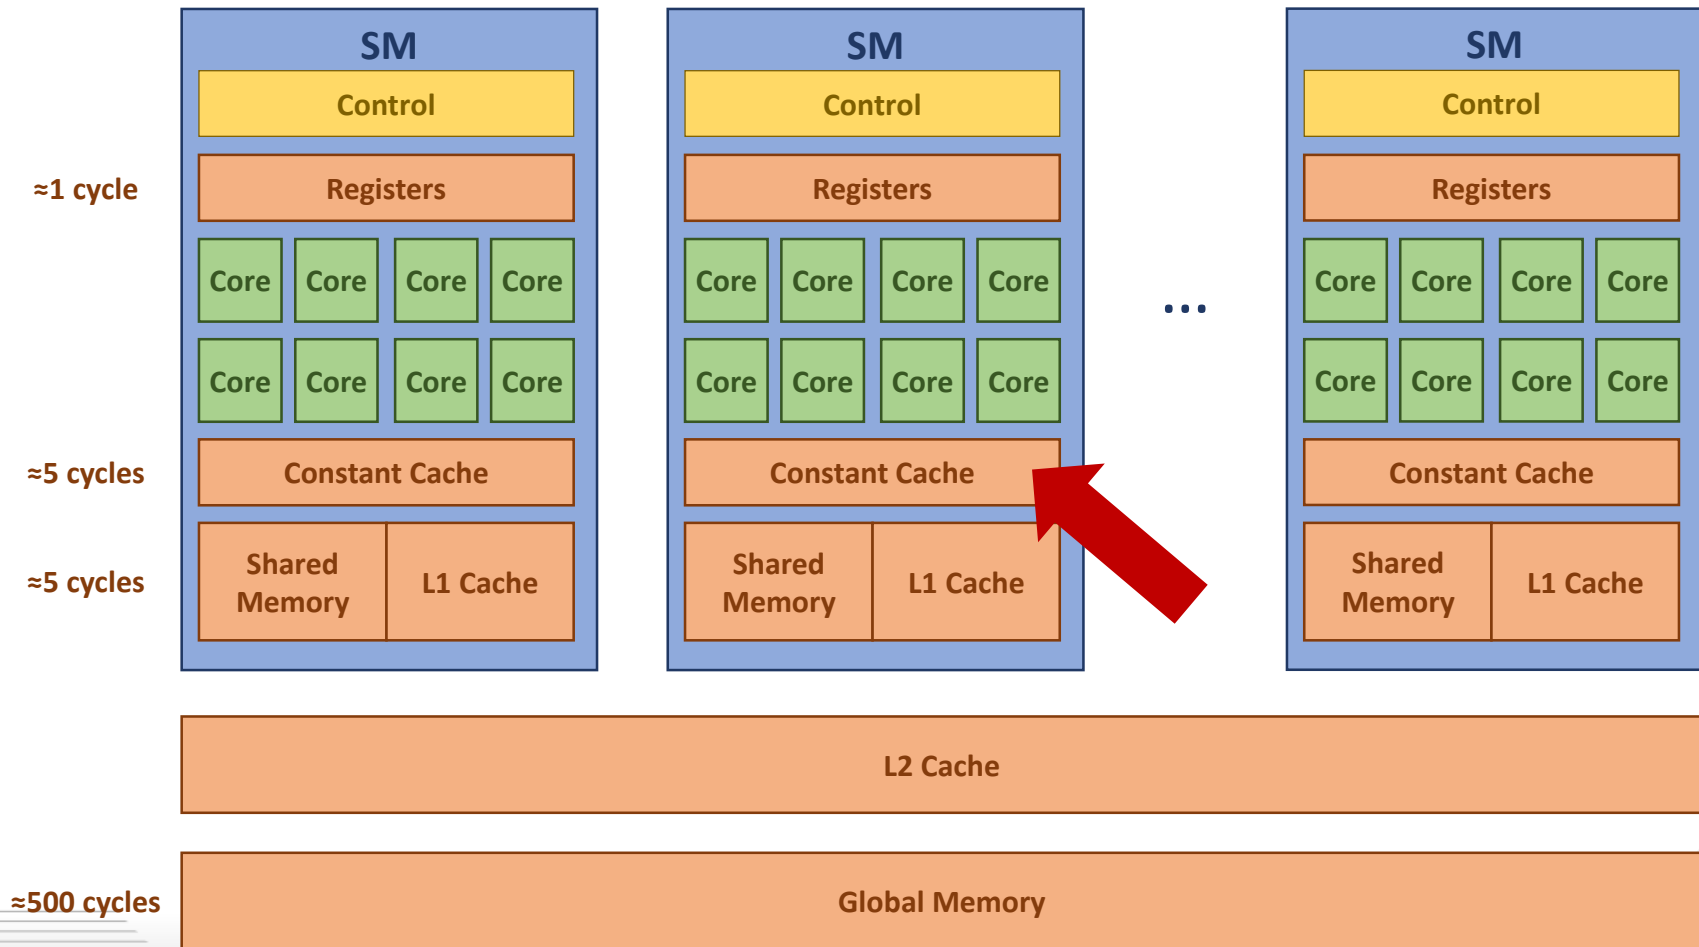
- Cannot modify during execution

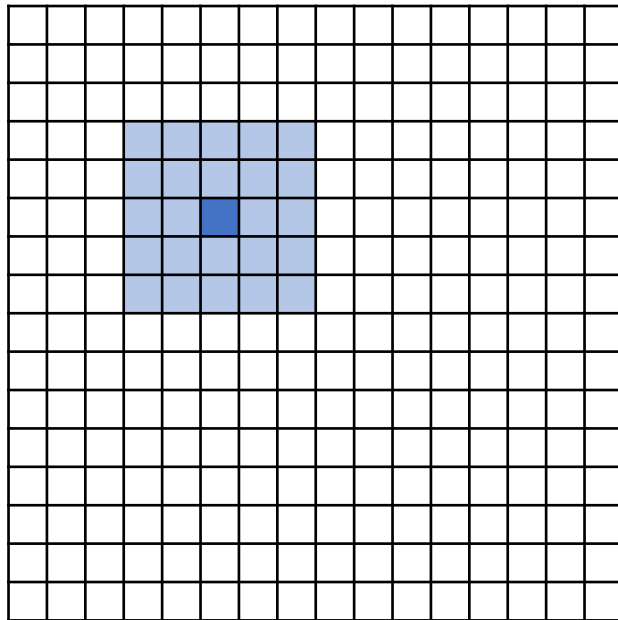
```
cudaMemcpyToSymbol(filter_c, filter, FILTER_DIM*FILTER_DIM*sizeof(float));
```

- Can only allocate up to 64KB
 - Otherwise, input is also constant, but it is too large to put in constant memory

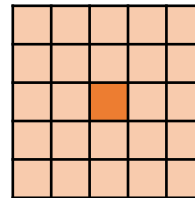
- Constant data: easier to build an efficient cache
 - No need to support write back
 - No need to support coherence
- Small size: minimize evictions
 - Cache for constant memory has low miss rate

Recall: Memory in the GPU Architecture

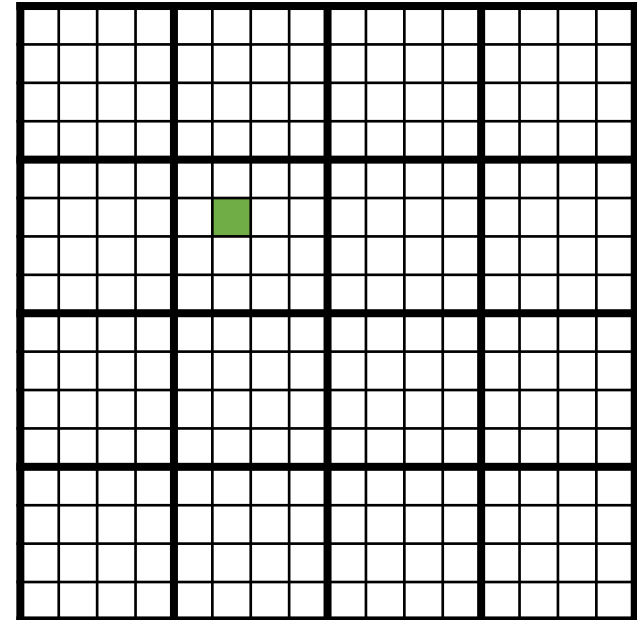




input
(in global
memory)



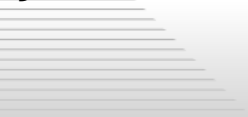
filter
(in constant
memory)

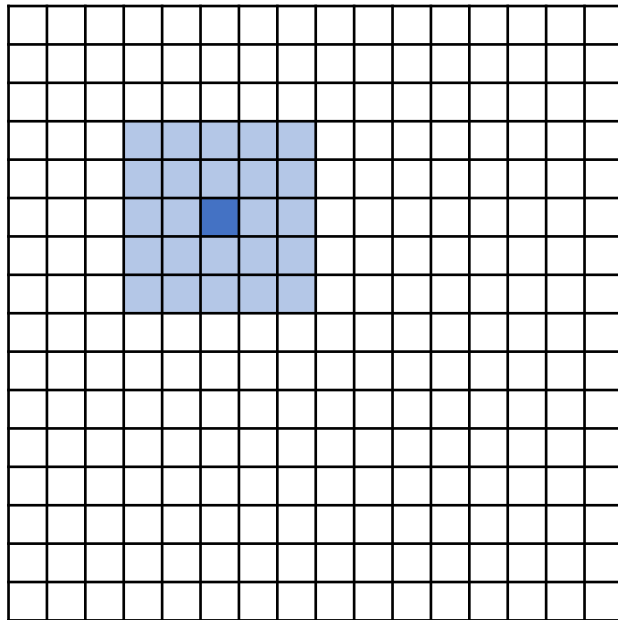


output

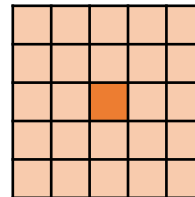
Parallelization approach: Assign one thread to compute each **output element** by looping over **input elements** and **filter** weights

```
__global__ void convolution_kernel(float* input, float* output, unsigned int width,
                                   unsigned int height) {
    int outRow = blockIdx.y*blockDim.y + threadIdx.y;
    int outCol = blockIdx.x*blockDim.x + threadIdx.x;
    if (outRow < height && outCol < width) {
        float sum = 0.0f;
        for(int filterRow = 0; filterRow < FILTER_DIM; ++filterRow) {
            for(int filterCol = 0; filterCol < FILTER_DIM; ++filterCol) {
                int inRow = outRow - FILTER_RADIUS + filterRow;
                int inCol = outCol - FILTER_RADIUS + filterCol;
                if(inRow >= 0 && inRow < height && inCol >= 0 && inCol < width) {
                    sum += filter_c[filterRow][filterCol]*input[inRow*width + inCol];
                }
            }
        }
        output[outRow*width + outCol] = sum;
    }
}
```

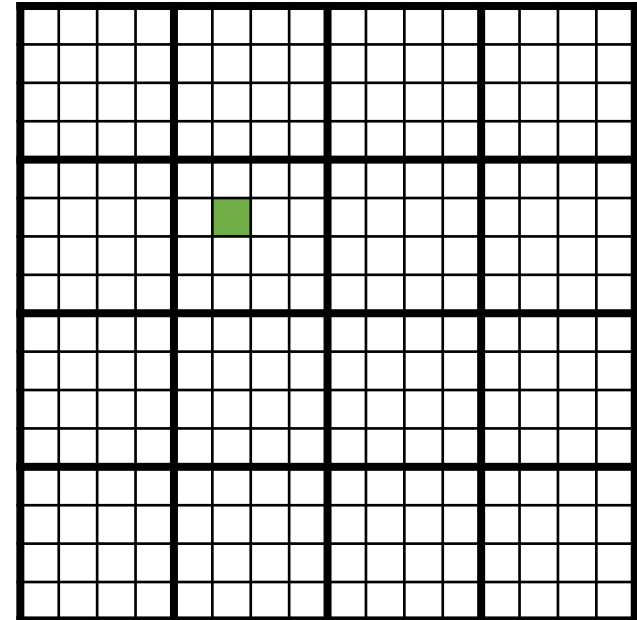




input
(in global
memory)

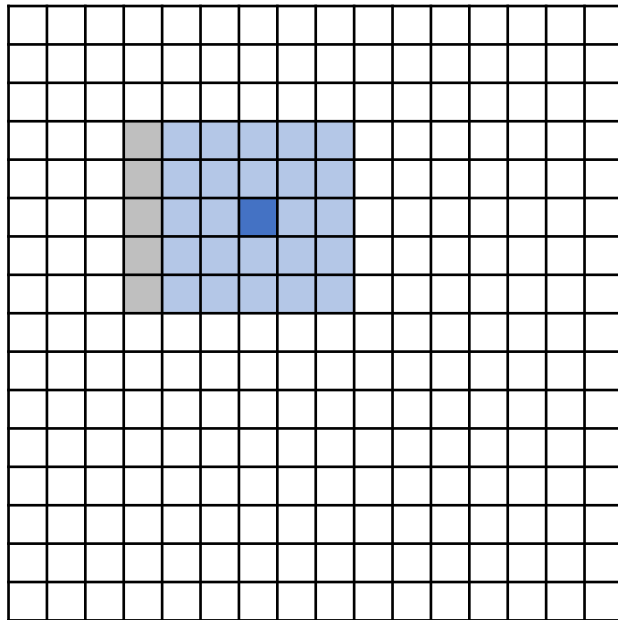


filter
(in constant
memory)

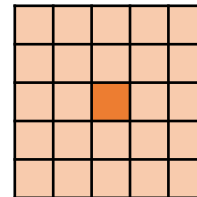


output

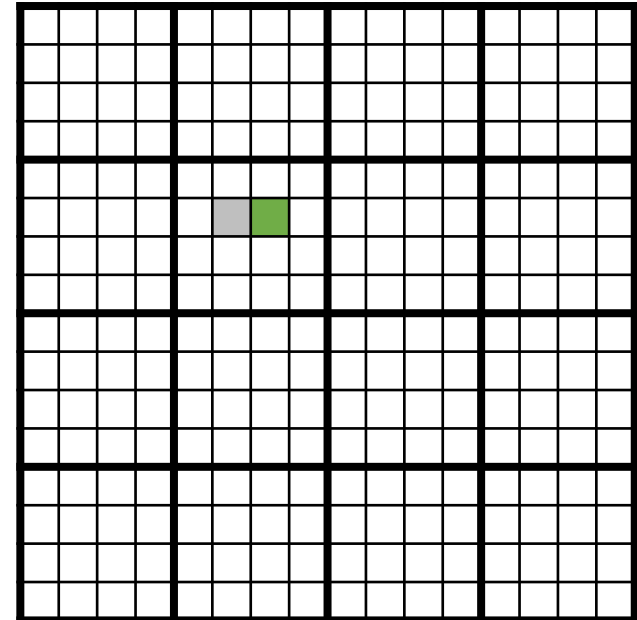
Observation: Threads in the same block load some of the same input elements



input
(in global
memory)

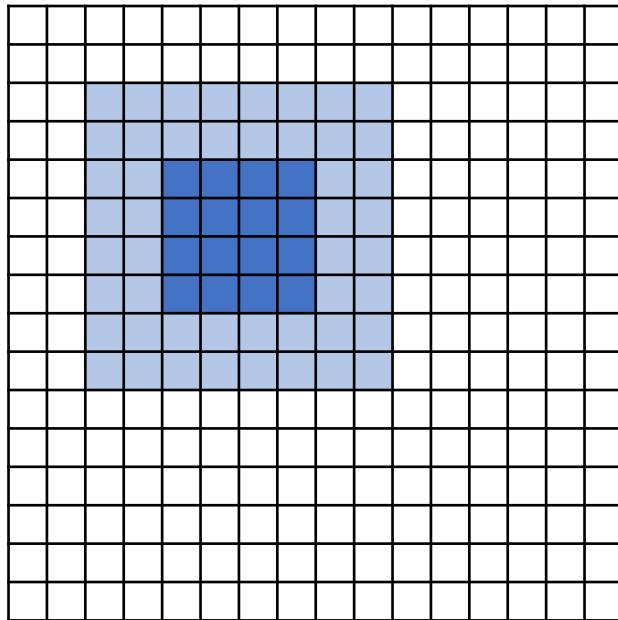


filter
(in constant
memory)

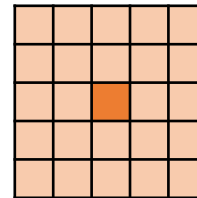


output

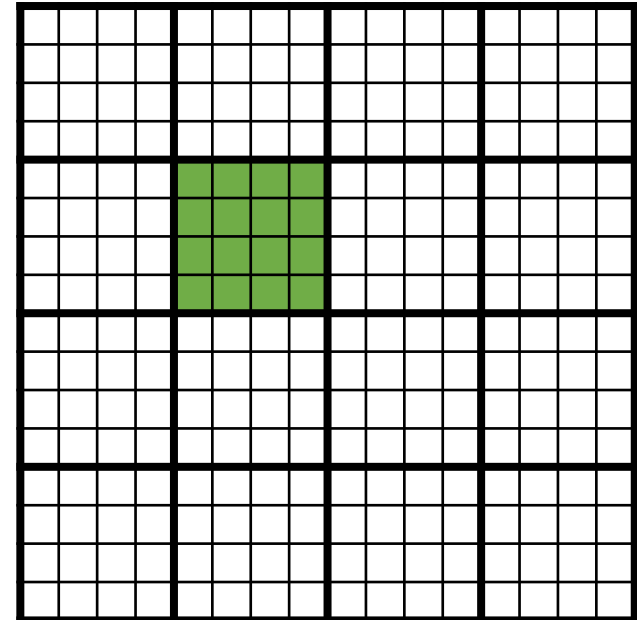
Observation: Threads in the same block load some of the same input elements



input
(in global
memory)



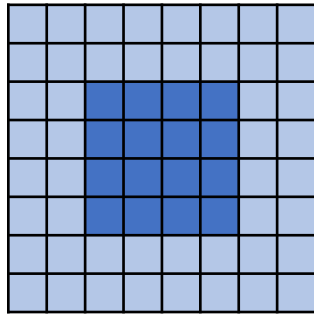
filter
(in constant
memory)



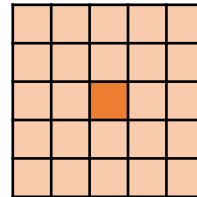
output

Observation: Threads in the same block load some of the same input elements

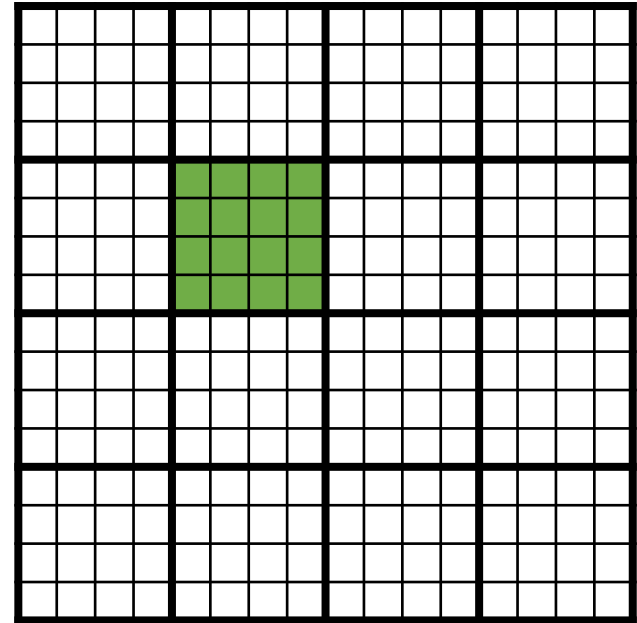
Optimization: Each thread loads one input element to shared memory and other threads access the element from shared memory



$\text{input}_{\text{tile}}$
(in shared
memory)

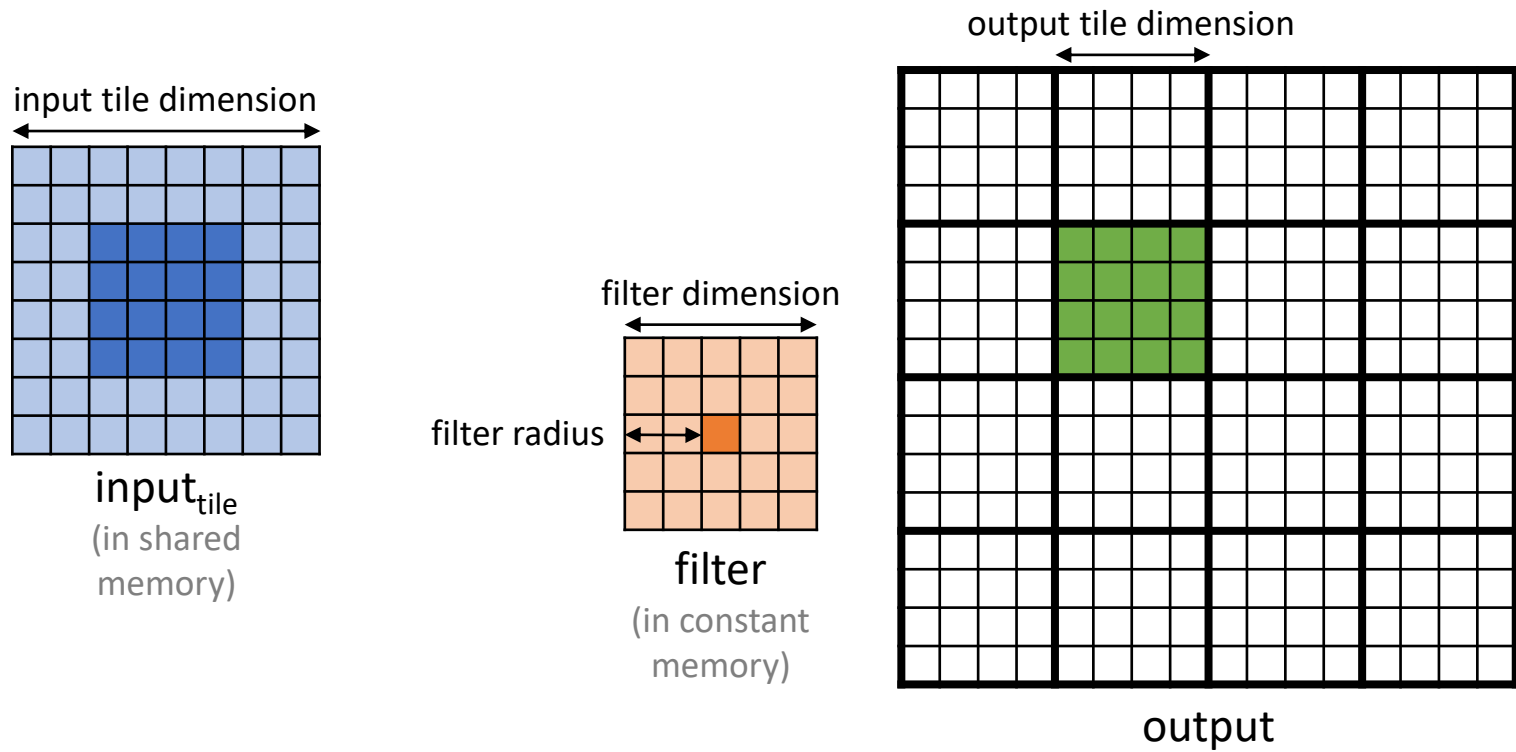


filter
(in constant
memory)



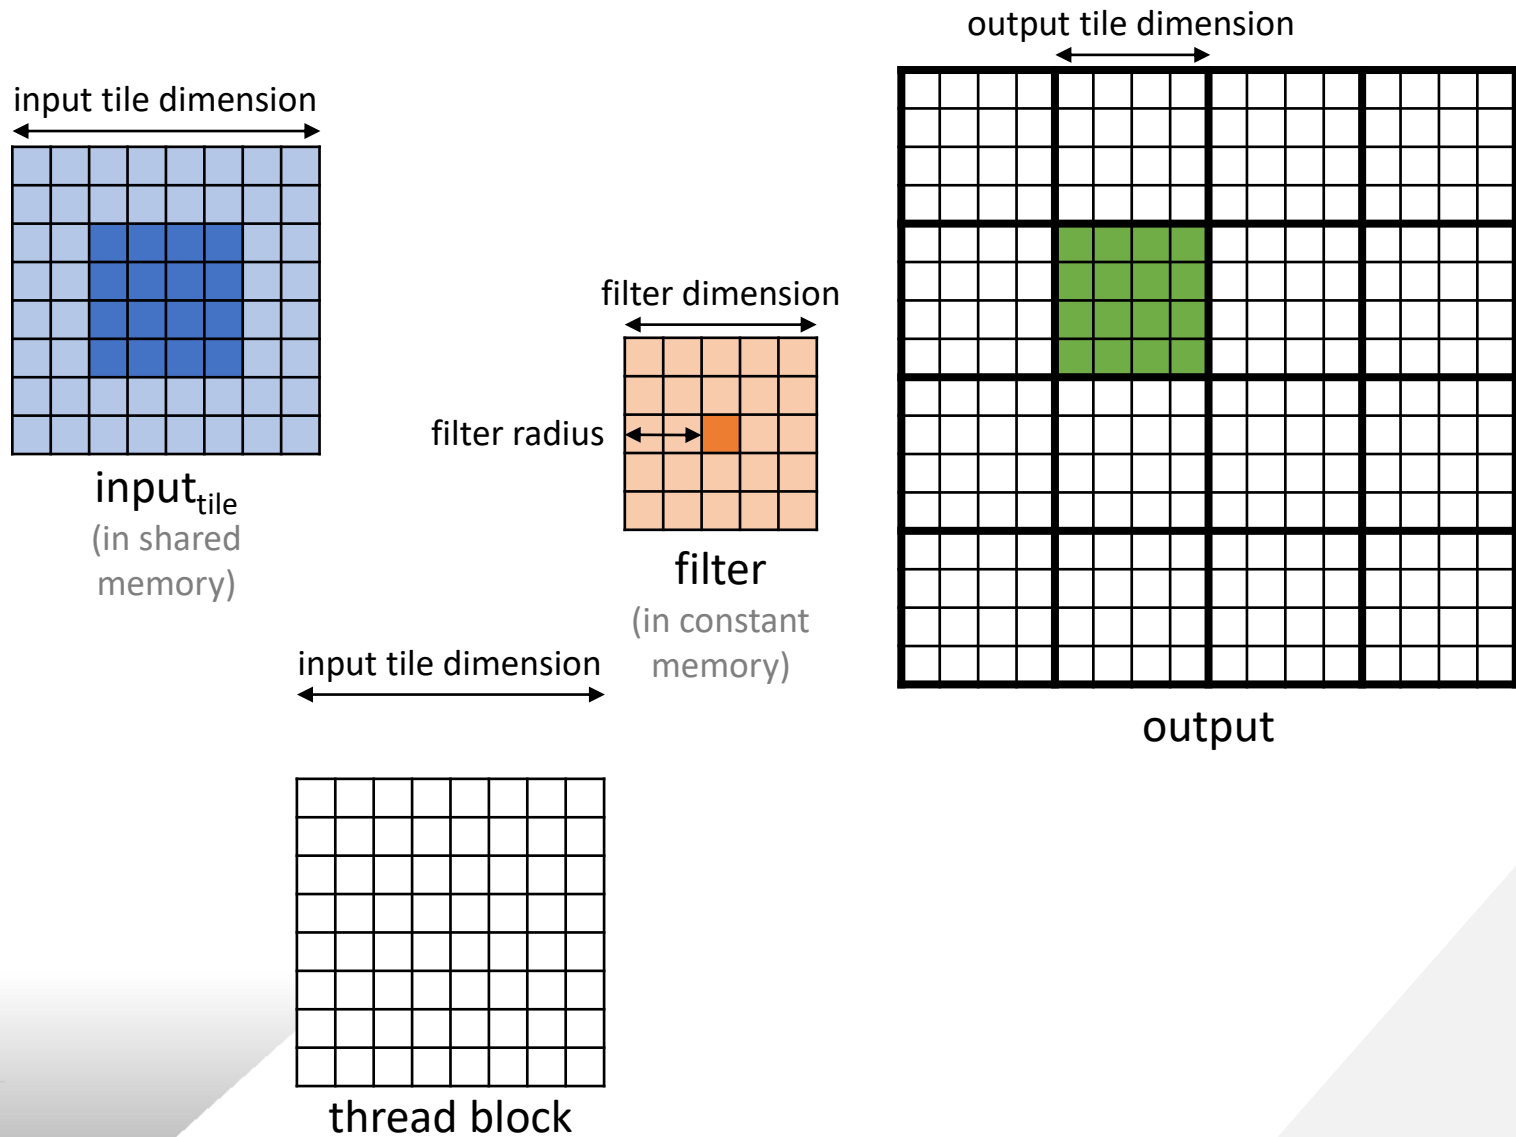
output

Optimization: Each thread loads one input element to shared memory and other threads access the element from shared memory

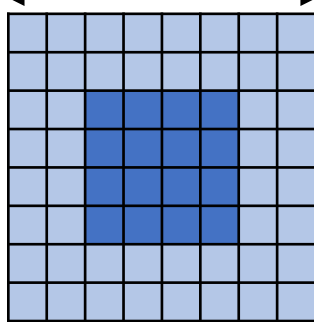


Challenge: Input and output tiles have different dimensions
(input tile dimension = output tile dimension + 2 × filter radius)

Solution: Launch enough threads per block to load the input tile to shared memory, then use a subset of them to compute and store the output tile

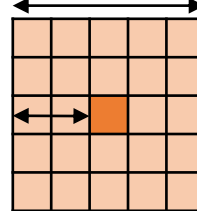


input tile dimension



$\text{input}_{\text{tile}}$
(in shared
memory)

filter dimension



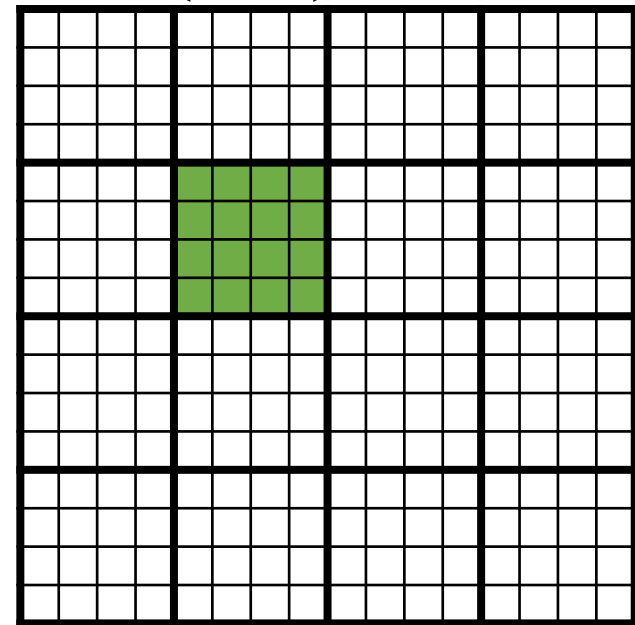
filter radius

filter

(in constant
memory)

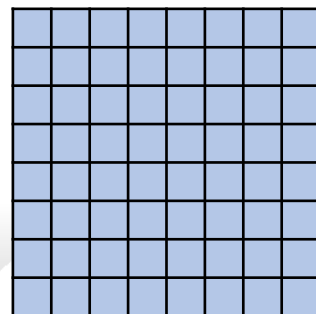
input tile dimension

output tile dimension



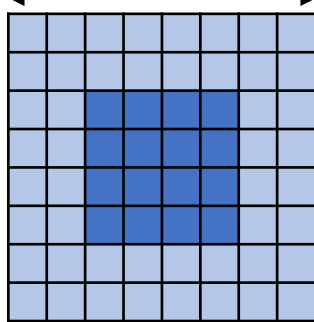
output

threads active when
loading input tile



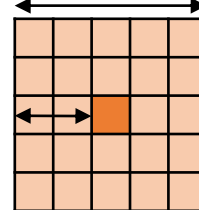
thread block

input tile dimension



$\text{input}_{\text{tile}}$
(in shared
memory)

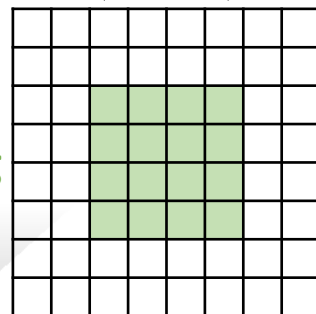
filter dimension



filter radius

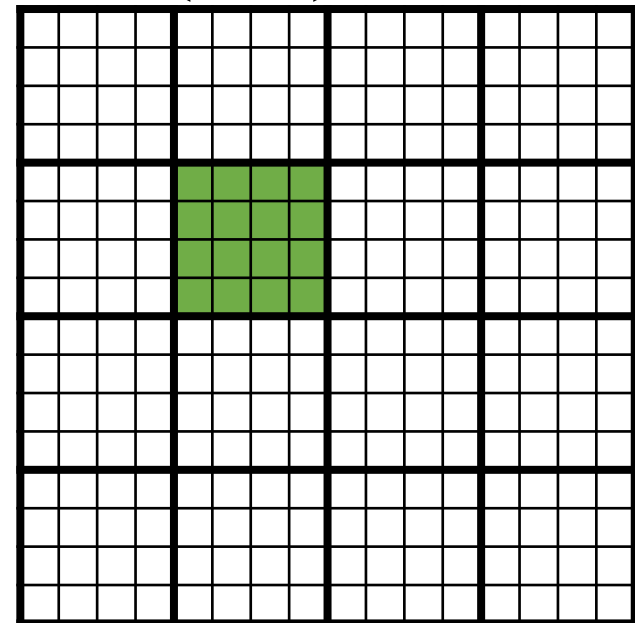
filter
(in constant
memory)

input tile dimension
output tile dimension



thread block

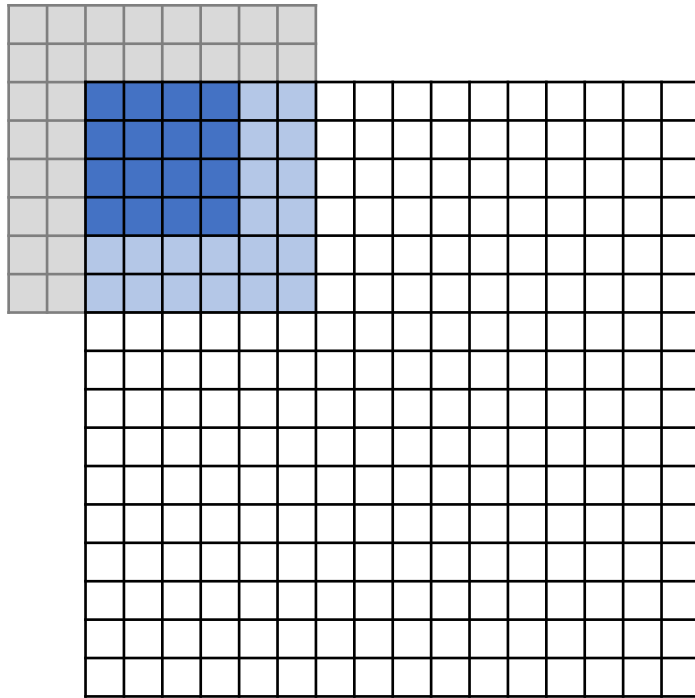
output tile dimension



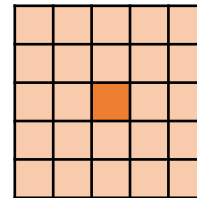
output

threads active when
computing and storing
the output tile

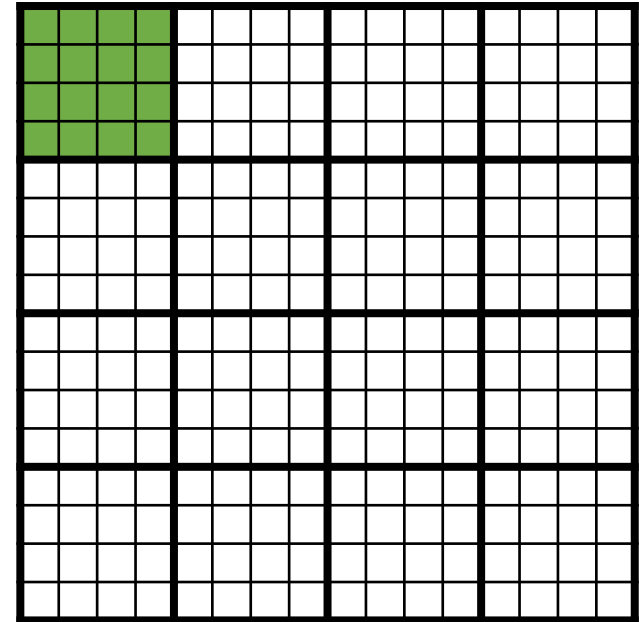
- Considering an $M \times M$ filter
- Without tiling:
 - Operations per thread: M^2 adds + M^2 muls = $2M^2$ OP
 - Global loads per thread: $M^2 \times 4 \text{ B} = 4M^2 \text{ B}$
 - Ratio: $(2M^2 \text{ OP}) / (4M^2 \text{ B}) = 0.5 \text{ OP/B}$
- With tiling:
 - Considering tile dimensions: input = T , output = $T-M+1$
 - Operations per block: $(T-M+1)^2 \times 2M^2 \text{ OP}$
 - Global loads per block: $T^2 \times 4 \text{ B}$
 - Ratio: $((T-M+1)^2 \times 2M^2 \text{ OP}) / (T^2 \times 4 \text{ B}) = 0.5M^2(1 - (M-1)/T)^2$
 - For $M=5$ and $T=32$: 9.57 OP/B ($\approx 19\times$ improvement!)



input
(in global
memory)



filter
(in constant
memory)



output

Threads computing output elements at the boundary access input elements that are out of bounds (also called *ghost* elements)

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0						
0	0						
0	0						
0	0						
0	0						
0	0						

$\text{input}_{\text{tile}}$
(in shared
memory)

filter
(in constant
memory)

output

Threads computing output elements at the boundary access input elements that are out of bounds (also called *ghost* elements)

Solution: Store zero to shared memory tile for our of bounds input elements

- Wen-mei W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022.