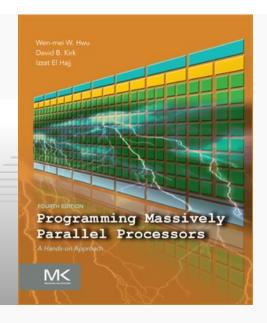


Programming Massively Parallel Processors

A Hands-on Approach

CHAPTER 11 > Prefix Sum (Scan)

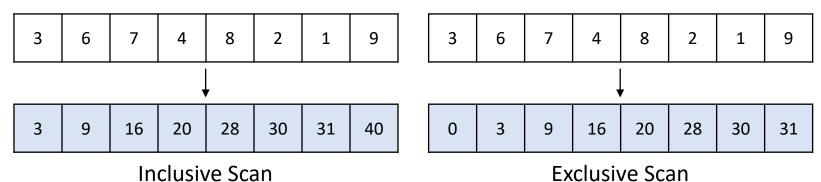




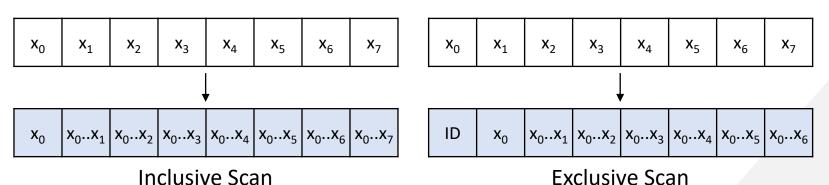
- A scan operation:
 - Takes:
 - An input array $[x_0, x_1, ..., x_{n-1}]$
 - An associative operator ⊕
 - e.g., sum, product, min, max
 - Returns:
 - An output array $[y_0, y_1, ..., y_{n-1}]$ where
 - Inclusive scan: $y_i = x_0 \oplus x_1 \oplus ... \oplus x_i$
 - Exclusive scan: $y_i = x_0 \oplus x_1 \oplus ... \oplus x_{i-1}$

Scan Example

• Addition example:



In general:



Sequential Scan

• Sequential scan for sum:

```
output[0] = input[0];
for(i = 1; i < N; ++i) {
    output[i] = output[i-1] + input[i];
}

Inclusive Scan

output[0] = 0.0f;
for(i = 1; i < N; ++i) {
    output[i] = output[i-1] + input[i-1];
}

Exclusive Scan</pre>
```

In general:

```
output[0] = input[0];
for(i = 1; i < N; ++i) {
    output[i] = f(output[i-1], input[i]);
}

Inclusive Scan

output[0] = IDENTITY;
for(i = 1; i < N; ++i) {
    output[i] = f(output[i-1], input[i-1]);
}

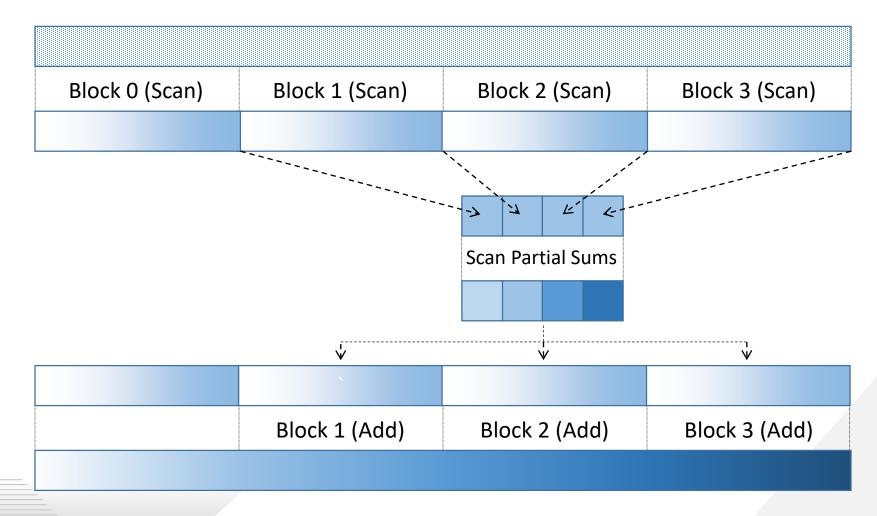
Exclusive Scan</pre>
```



Segmented Scan

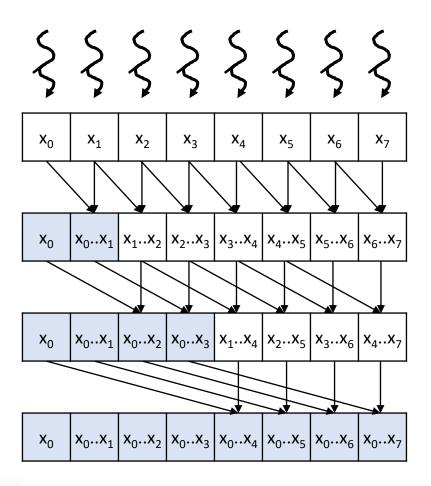
- Similar to reduction, threads must synchronize to perform scan
 - Cannot synchronize across threads in different blocks
- Solution: segmented scan (or hierarchical scan)
 - Every thread block scans a segment
 - Scan the partial sums
 - Update segments based on partial sums

Segmented Scan Example



Still need to implement a parallel in each block





Intuition: equivalent to having a reduction tree for each element, with overlap between trees

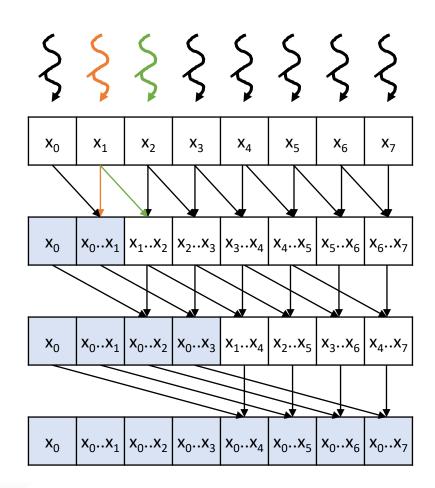
Kogge-Stone Parallel (Inclusive) Scan Code

```
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
output[i] = input[i];
__syncthreads();
for(unsigned int stride = 1; stride <= BLOCK_DIM/2; stride *= 2) {</pre>
    if(threadIdx.x >= stride) {
        output[i] += output[i - stride];
    __syncthreads();
}
                                                        Incorrect!
                                                   Different threads are
if(threadIdx.x == BLOCK_DIM - 1) {
                                                  reading and writing the
    partialSums[blockIdx.x] = output[i];
}
                                                same data location without
                                                       synchronizing
```



Thread 1 may update value at index 1 before thread 2 reads it


Solution: wait for everyone to read before updating



Kogge-Stone Parallel (Inclusive) Scan Code

```
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
output[i] = input[i];
__syncthreads();
for(unsigned int stride = 1; stride <= BLOCK_DIM/2; stride *= 2) {</pre>
    float v;
    if(threadIdx.x >= stride) {
        v = output[i - stride];
                                        Wait for everyone to read
    __syncthreads();
                                             before updating
    if(threadIdx.x >= stride) {
        output[i] += v;
    __syncthreads();
}
if(threadIdx.x == BLOCK_DIM - 1) {
    partialSums[blockIdx.x] = output[i];
}
```

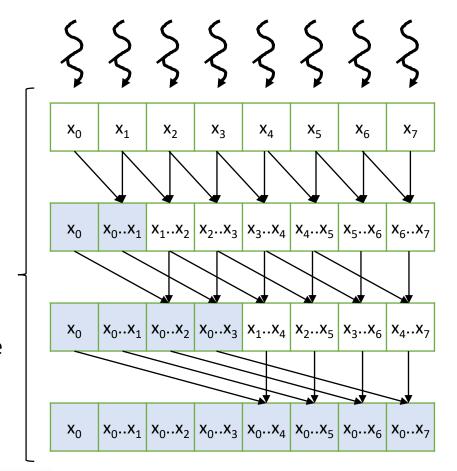




Observation: memory locations are reused



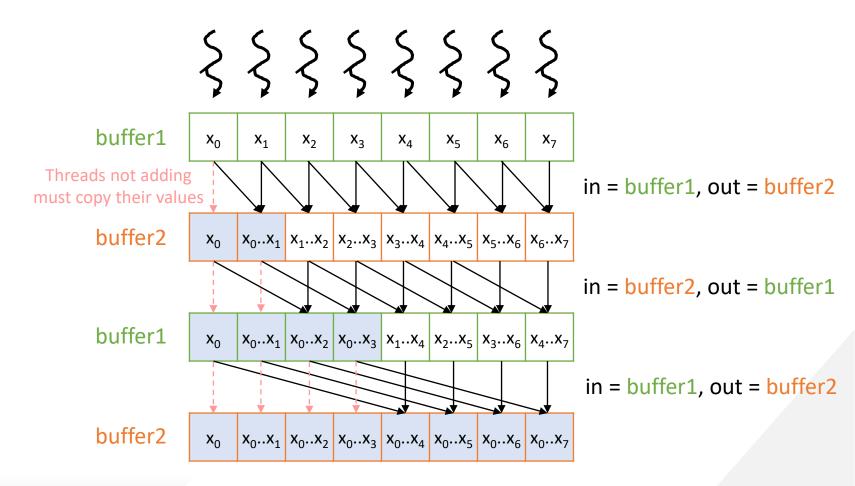
Optimization: load once to a shared memory buffer and perform successive reads and writes to the same array can be done in shared memory



Using Shared Memory Code

```
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
__shared__ float buffer_s[BLOCK_DIM];
buffer_s[threadIdx.x] = input[i];
__syncthreads();
for(unsigned int stride = 1; stride <= BLOCK_DIM/2; stride *= 2) {</pre>
    float v:
    if(threadIdx.x >= stride) {
        v = buffer_s[threadIdx.x - stride];
    _syncthreads();
    if(threadIdx.x >= stride) {
        buffer_s[threadIdx.x] += v;
    __syncthreads();
}
if(threadIdx.x == BLOCK_DIM - 1) {
    partialSums[blockIdx.x] = buffer_s[threadIdx.x];
}
output[i] = buffer_s[threadIdx.x];
```





<u>Optimization:</u> eliminate one synchronization by using two buffers and alternating them as the input/output buffer (called double buffering)

Double Buffering

```
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
__shared__ float buffer1_s[BLOCK_DIM];
<u>__shared__</u> float buffer2_s[BLOCK_DIM];
float* inBuffer_s = buffer1_s;
float* outBuffer_s = buffer2_s;
inBuffer_s[threadIdx.x] = input[i];
__syncthreads();
for(unsigned int stride = 1; stride <= BLOCK_DIM/2; stride *= 2) {</pre>
    if(threadIdx.x >= stride) {
        outBuffer_s[threadIdx.x] =
                inBuffer_s[threadIdx.x] + inBuffer_s[threadIdx.x - stride];
    } else {
        outBuffer_s[threadIdx.x] = inBuffer_s[threadIdx.x];
    __syncthreads();
    float* tmp = inBuffer_s;
    inBuffer_s = outBuffer_s;
    outBuffer_s = tmp;
}
if(threadIdx.x == BLOCK_DIM - 1) {
    partialSums[blockIdx.x] = inBuffer_s[threadIdx.x];
output[i] = inBuffer_s[threadIdx.x];
```



Exclusive Scan

- Formulate as inclusive scan:
 - Shift elements by one when loading them from shared memory (skips the last element)
 - Fetch last element when computing partial sum

Exclusive Scan

```
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
__shared__ float buffer1_s[BLOCK_DIM];
__shared__ float buffer2_s[BLOCK_DIM];
float* inBuffer_s = buffer1_s;
float* outBuffer_s = buffer2_s;
if(threadIdx.x == 0) {
    inBuffer_s[threadIdx.x] = 0.0f;
} else {
    inBuffer_s[threadIdx.x] = input[i - 1];
__syncthreads();
for(unsigned int stride = 1; stride <= BLOCK_DIM/2; stride *= 2) {</pre>
    if(threadIdx.x >= stride) {
        outBuffer_s[threadIdx.x] =
                inBuffer_s[threadIdx.x] + inBuffer_s[threadIdx.x - stride];
    } else {
        outBuffer_s[threadIdx.x] = inBuffer_s[threadIdx.x];
    __syncthreads();
    float* tmp = inBuffer_s;
    inBuffer_s = outBuffer_s;
    outBuffer_s = tmp;
}
if(threadIdx.x == BLOCK_DIM - 1) {
    partialSums[blockIdx.x] = inBuffer_s[threadIdx.x] + input[i];
output[i] = inBuffer_s[threadIdx.x];
```

Work Efficiency

- A parallel algorithm is work-efficient if it performs the same amount of work as the corresponding sequential algorithm
- Scan work efficiency
 - Sequential scan performs N additions
 - Kogge-Stone parallel scan performs:
 - log(N) steps, N 2^{step} operations per step
 - Total: (N-1) + (N-2) + (N-4) + ... + (N-N/2)= N*log(N) - (N-1) = O(N*log(N)) operations
 - Algorithm is not work efficient
- If resources are limited, parallel algorithm will be slow because of low work efficiency

