



**BITS Pilani**  
Pilani Campus

# Lock Free Concurrent Data Structures

K Hari Babu  
Department of Computer Science & Information Systems

# Problems with Locking



- Performance issues
  - Locks necessitate waits, context switching, CPU stalls, etc... effecting performance
- Undoing locks
  - If a thread dies while holding a lock, no other thread can proceed
- Can lead to deadlocks
- Priority inversion
  - If a higher priority thread needs to enter critical region but lock is held by low priority thread.

# Lock-free Programming



- Thread-safe access to shared data without the use of synchronization primitives such as mutexes
- Possible but not practical in the absence of hardware support

# General Approach to Lock-Free Algorithms



- Designing generalized lock-free algorithms is hard
- Design lock-free data structures instead
  - buffer, list, stack, queue, map etc.
- One can't implement lock free algorithms in terms of lock-based data structures.

# Lock-free Data Structures



- Is it possible to build data structures that are thread-safe without locks?
  - Yes
- Lock-free data structures
  - Include no locks, but are thread safe
  - may introduce starvation
    - Due to retry loops
- Wait-free data structures
  - Include no locks, are thread safe, and avoid starvation
  - Wait-free implies lock-free. Wait-free is much stronger than lock-free
  - Wait-free structures are very hard to implement
  - Impossible to implement for many data structures
  - Often restricted to a fixed number of threads

# Lock-free Data Structures



- Very few standard libraries/APIs implement these data structures
- Implementations are often platform-dependent
  - Rely on low-level assembly instructions
  - Many structures are very new, not widely known
- Lock free data structures make use of the retry loop pattern
  - Read some state
  - Do a useful operation
  - Attempt to modify global state if it hasn't changed (using Compare and Swap (CAS))
- This is similar to a spinlock
  - But, the assumption is that wait times will be small
  - However, retry loops may introduce starvation

# Instruction-level Atomicity



- On x86
  - The INC instruction adds 1 to the destination operand. The destination operand can be a register or a memory location.
  - In case of memory address, this instruction loads the value, increments and stores it back. Meanwhile DMA or other cores may access the same memory location, and modify it.
- x86 provides LOCK prefix which ensures that the CPU has exclusive ownership of the appropriate cache line for the duration of the operation and by asserting a bus lock
- The lock prefix makes an instruction atomic and is legal only with some instructions.

# Compare and Swap Instruction (CAS)



- On x86, known as compare and exchange

```
2  spin_lock:
3      mov ecx, 1
4      mov eax, 0
5      lock cmpxchg ecx, [lock_addr]
6      jnz spinlock
```

- cmpxchg compares eax and the value of lock\_addr
- If eax == [lock\_addr], swap ecx and [lock\_addr]



# The Price of Atomicity



- Atomic operations are very expensive on a multi-core system
  - Caches must be flushed
    - CPU cores may see different values for the same variable if they have out-of-date caches
    - Cache flush can be forced using a memory fence (sometimes called a memory barrier). MFENCE performs a serializing operation on all load-from-memory and store-to-memory instructions that were issued prior to the MFENCE instruction
  - Memory bus must be locked
    - No concurrent reading or writing
  - Other CPUs may stall
    - May block on the memory bus or atomic instructions

# Implementation



- Many lock-free data structures can be built using compare and swap (CAS)

```
2 bool cas(int * addr, int oldval, int newval) {  
3     if (*addr == oldval) {  
4         *addr = newval;  
5         return true;  
6     }  
7     return false;  
8 }
```

- This can be done atomically on x86 using the **cmpxchg** instruction
- Many compilers have built in atomic swap functions
  - GCC: `__sync_bool_compare_and_swap(ptr, oldval, newval)`
  - MSVC: `InterlockedCompareExchange(ptr,oldval,newval)`

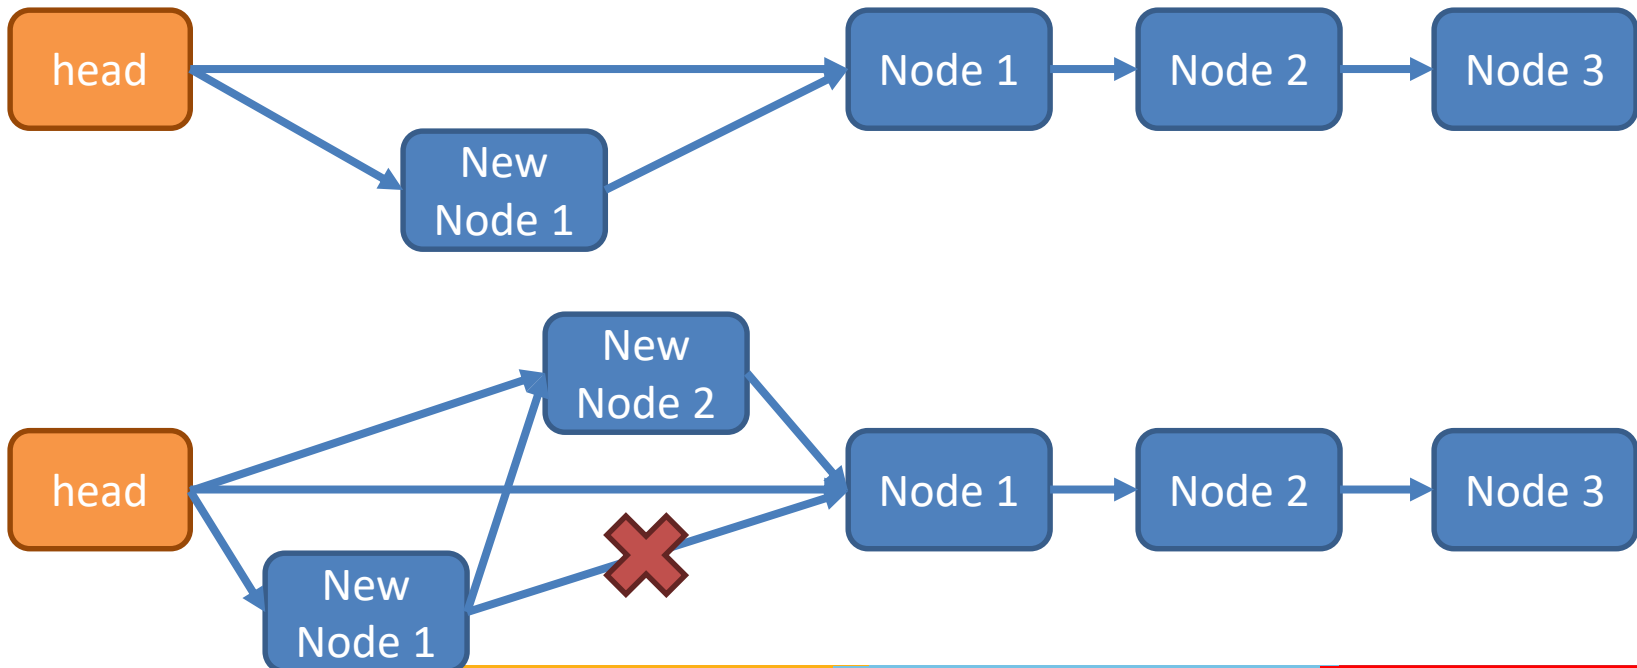
# Lock-free Stack Example: Push



```
1 class Node {  
2     Node * next;  
3     int data;  
4 };  
5  
6 // Root of the stack  
7 volatile Node * head;
```

```
1 void push(int t) {  
2     Node* node = new Node(t);  
3     do {  
4         node->next = head->next;  
5     } while (!cas(head->next, node->next, node));  
6 }
```

```
2 bool cas(int * addr, int oldval, int newval) {  
3     if (*addr == oldval) {  
4         *addr = newval;  
5         return true;  
6     }
```



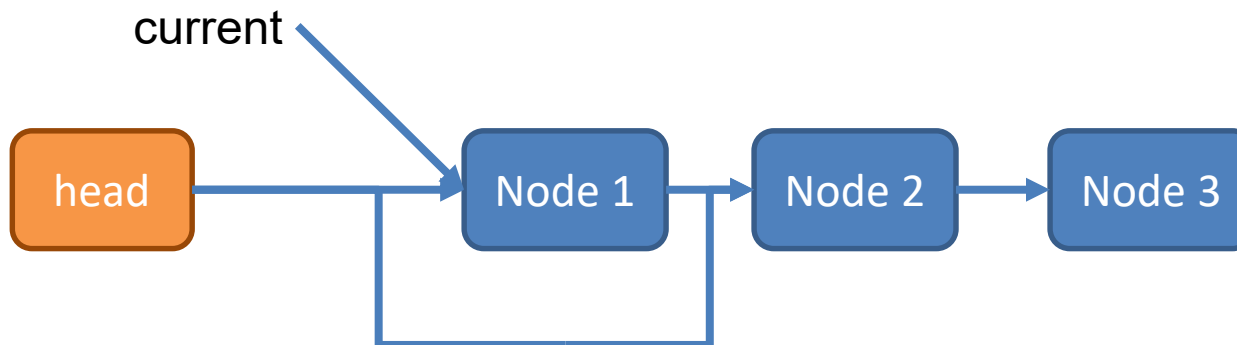
# Lock-free Stack Example: Pop



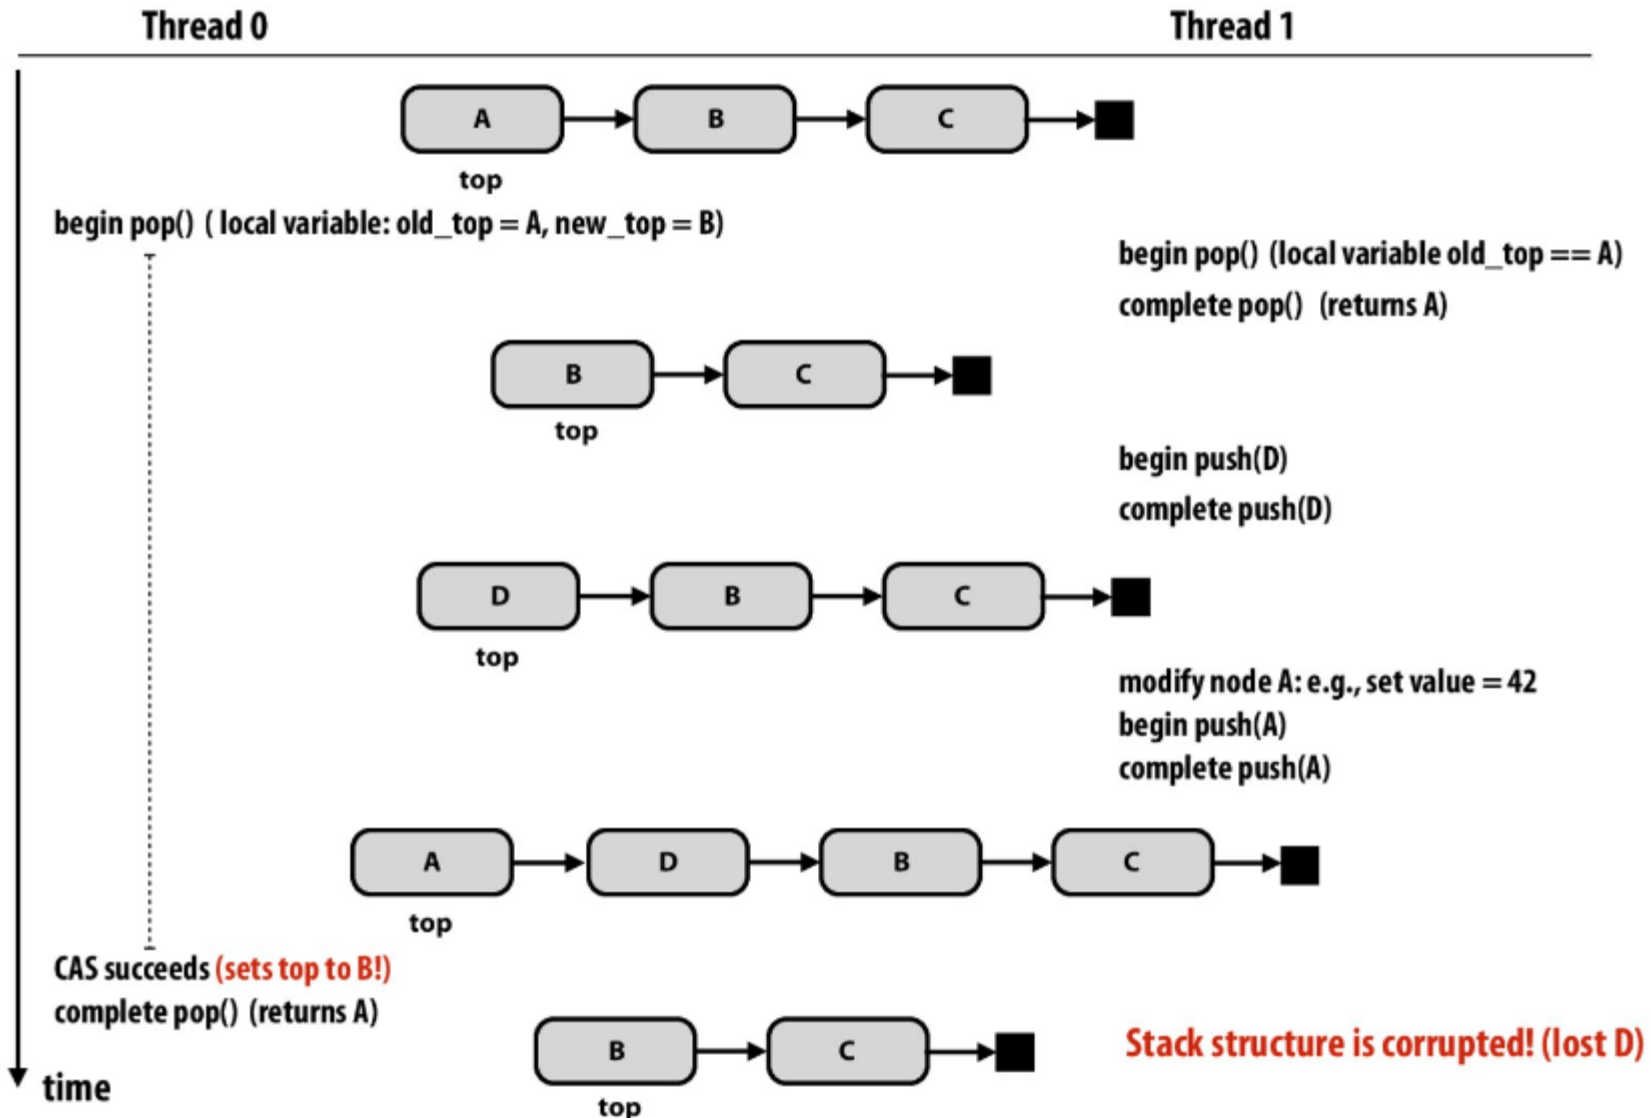
```
1 class Node {  
2     Node * next;  
3     int data;  
4 };  
5  
6 // Root of the stack  
7 volatile Node * head;
```

```
1 bool pop(int& t) {  
2     Node* current = head->next;  
3     while(current) {  
4         if(cas(head->next, current, current->next)) {  
5             t = current->data;  
6             delete current;  
7             return true;  
8         }  
9         current = head->next;  
10    }  
11    return false;  
12 }
```

```
2 bool cas(int * addr, int oldval, int newval) {  
3     if (*addr == oldval) {  
4         *addr = newval;  
5         return true;  
6     }  
7     return false;  
8 }
```



# Lock-free Stack: ABA problem



- Workarounds
  - Keep a 'update count' (needs 'doubleword CAS')
    - Kind of version counter that will not match though value matches
  - Don't recycle the memory 'too soon'

# Applications of Lock-Free Structures



- Stack
- Queue
- Linked list
- Doubly linked list
- Hash table
- Many variations on each
  - Lock free vs. wait free
- Memory managers
  - Lock free malloc() and free()
- The Linux kernel
  - Many key structures are lock-free

# References



- Geoff Langdale, Lock-free Programming

[http://www.cs.cmu.edu/~410-s05/lectures/L31\\_LockFree.pdf](http://www.cs.cmu.edu/~410-s05/lectures/L31_LockFree.pdf)



# Q&A





**BITS Pilani**  
Pilani Campus



**Thank You**