



BITS Pilani
Pilani Campus

PRAM Model

K Hari Babu (Slides adapted from Prof. Shan's)
Department of Computer Science & Information Systems



BITS Pilani
Pilani Campus



RAM Model

Why Machine Models?

- What is a machine model?
 - An abstraction that describes the operation of a machine
 - Associates a value (cost) with each machine operation
- Why do we need models?
 - Makes it easier to analyze and develop algorithms
 - Hides the machine implementation details so that general results that apply to a broad class of machines are obtainable
 - Analyzes the achievable complexity (time, space, etc.) bounds
 - Analyzes maximum parallelism
 - Conversely, models are directly related to algorithms.

RAM (Random Access Machine) Model

- Memory consists of infinite array (memory cells).
- Instructions executed sequentially, one at a time
- All instructions take unit time:
 - Load/store
 - Arithmetic
 - Logic
- Running time of an algorithm: the number of instructions executed
- Memory requirement: the number of memory cells used in the algorithm

RAM (Random Access Machine) Model

- The RAM model is the base of algorithm analysis for sequential algorithms although it is not perfect:
 - Memory is not infinite
 - Not all memory accesses take the same time
 - Not all arithmetic operations take the same time
 - Instruction pipelining is not taken into consideration
- The RAM model (with asymptotic analysis) often gives relatively realistic results

PRAM (Parallel RAM)

- An unbounded collection of processors
- Each process has an infinite number of registers
- An unbounded collection of shared memory cells
- All processors can access all memory cells in unit time (when there is no memory conflict)
- All processors execute PRAM instructions synchronously (some processors may be idle)
- Each PRAM instruction executes in a 3-phase cycle:
 - Read from a share memory cell (if needed)
 - Computation
 - Write to a share memory cell (if needed)
- PRAM is an abstract machine model
 - multiple execution units, single control and clock, local (private) memory units, and global (shared) memory.

PRAM (Parallel RAM)

- The only way processors exchange data is through the shared memory.
- Individual processor
 - Time: number of instructions executed
 - Space: number of memory cells accessed
- Parallel time complexity:
 - the number of synchronous steps in the algorithm
 - Time taken by the longest running process
- Space complexity:
 - the number of shared memory cells
- Parallelism:
 - the number of processors used

PRAM Model - variants

- Shared Memory Model:
 - Can processes – *running in different processors* – access locations in global shared memory concurrently?
 - If yes, is concurrent access allowed for **read** as well as for **write** operations?
 - If yes, how are concurrent writes governed?
- Variants of PRAM are defined on the basis of answers to these questions.

PRAM Model - variants

- Variants based on concurrent operations:
- Exclusive Read Exclusive Write (EREW)
 - Shared memory locations cannot be read / written concurrently.
- Concurrent Read Exclusive Write (CREW)
 - Shared memory locations can be read but not written concurrently.
- Concurrent Read Concurrent Write (CRCW)
 - Shared memory locations can be read / written concurrently.

PRAM variants: Interpretation at software level

- EREW model

- Contract between the processor and the program:
 - processor does not support concurrent memory operations (read or write);
 - program must devise its own mechanism for ensuring memory is accessed serially.

- CREW model

- Contract between the processor and the program:
 - processor supports concurrent reads;
 - program must devise its own mechanism for concurrent writes to get processed serially.

PRAM Model: CRCW - variants

- Variants of CRCW based on handling of write conflicts:
 - Shared memory locations can be read / written concurrently:
 - Common: All concurrent writes must write the same value
 - Arbitrary: Concurrent writes may write different values but which of those values gets stored is arbitrary.
 - Priority: Concurrent writes may write different values and the value that gets written stored is decided by priority
 - priority is the id of the writer i.e. the processor

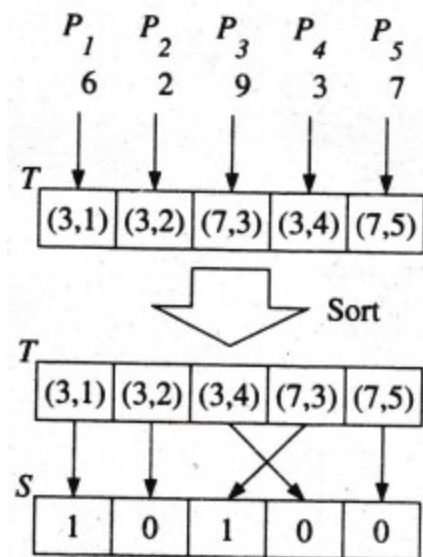
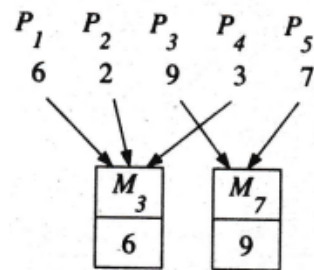
Relative Strengths of Models

- Model A is computationally stronger than model B if and only if any algorithm written in B will run unchanged in A. We can prove,
 - $EREW \leq CREW \leq CRCW \text{ (common)} \leq CRCW \text{ (arbitrary)} \leq CRCW \text{ (Priority)}$
 - Weakest model Strongest Model
 - Most realisticLeast Realistic
 - CREW PRAM can execute any EREW PRAM algorithm in the same amount of time. Concurrent read facility is not used.
 - CRCW PRAM can execute any EREW PRAM algorithm
 - Priority PRAM model is the strongest
 - An algorithm to solve a problem on the EREW PRAM model can have higher time complexity than an algorithm on Priority model

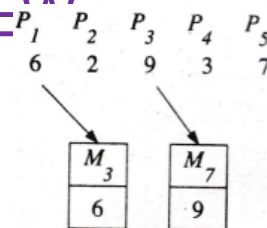
Priority PRAM vs EREW PRAM

- A p-processor EREW PRAM can sort a p-element array stored in global memory in $\Theta(\log p)$ time
- Theorem: A p-processor Priority PRAM can be simulated by a p-processor EREW PRAM with the time complexity increased by a factor of $\Theta(\log p)$

Handling of concurrent write in Priority PRAM



Simulating concurrent write on EREW



Priority PRAM vs EREW PRAM

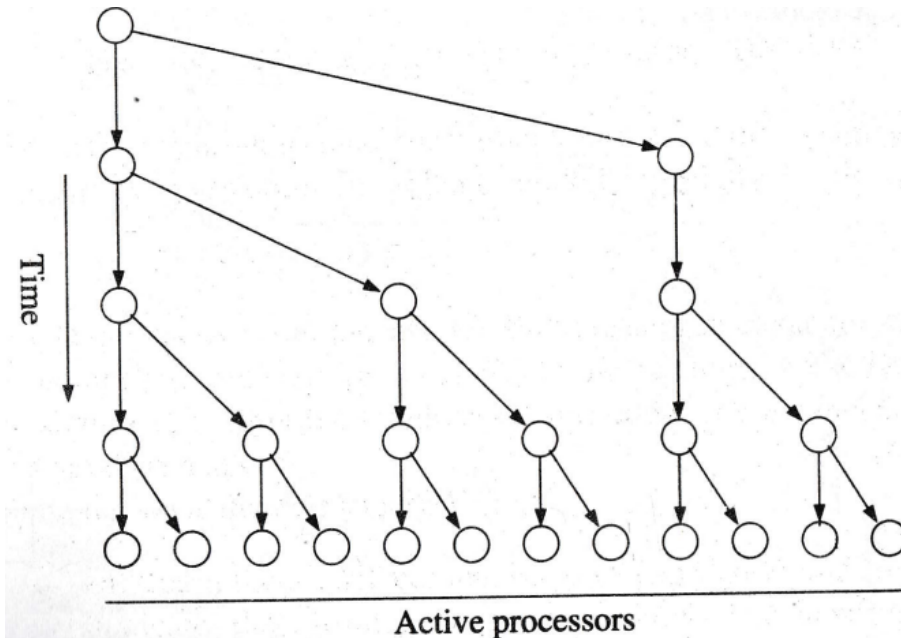
- A p-processor EREW PRAM can sort a p-element array stored in global memory in $\Theta(\log p)$ time
- Theorem: A p-processor Priority PRAM can be simulated by a p-processor EREW PRAM with the time complexity increased by a factor of $\Theta(\log p)$
 - Proof:
 - All write operations at a time by the Priority PRAM are stored in an array T of pairs $(M[j], P_i)$.
 - Sort T in lexicographic order
 - $\Theta(\log p)$ time. (parallel sorting with p processors)
 - Find the highest priority processor writing into a specific location
 - $\Theta(1)$ time

PRAM: Algorithmic Model

- Time Complexity vs. Cost
 - $\text{Cost} = \text{Time Complexity} * \text{\#processors}$
- Since a PRAM algorithm begins only with a single processor active, PRAM algorithms have two phases
- 2-phase model:
 - Activation Phase and Computation Phase
 - In the first phase sufficient number of processors are activated (spawned)
 - In the second phase these activated processors perform the computation in parallel

Activation Phase

- Exactly $\text{ceil}(\log p)$ processor steps are necessary and sufficient to change from 1 active processor to p active processors
 - Because the number of active processors can double by executing a single instruction



All logarithms are base 2 unless stated

Computation Phase

- After the necessary processes are spawned, each processor executes same segment of code in parallel

```
for all <processor list>  
do  
  <statement list>  
endfor
```

- This is essentially similar to programming on SIMD architectures:
 - also referred to as the SPMD model

PRAM: Algorithmic Model - Computation

- If all processors execute the same set of statements, how do you obtain variations?
 - e.g. different data to be processed in each processor
 - e.g. different choices to be made in each processor
- Use the processor id (rank) for data access or control:

```
for all Pi where i = 0 to p-1
do
    x = M[i];
    if (i%2 == 0)
        { M[i] = x * x; }
    else
        { M[i] = x * x * x; }
endfor
```

PRAM Algorithms – Example 1

- Algorithm for Adding Two Matrices [Mat-Add] :
 - How many processors?
 - Consider the two algorithms below.
- Algorithm 1 [for Mat-Add]:
 - // Input: Matrices A and B of size $m \times n$
 - for** all $P_{i,j}$ in $i = 1$ to m , $j = 1$ to n
 - {** $C[i,j] = A[i,j] + B[i,j]$ **}**
 - $m \times n$ processors, $O(1)$ time, Cost = $O(m \times n)$
- Algorithm 2 [for Mat-Add]:
 - for** all P_i in $i = 1$ to m {
 - for** $j = 1$ to n { $C[i,j] = A[i,j] + B[i,j]$ **}}**
 - m processors, $O(n)$ time, Cost = $O(m \times n)$

PRAM Algorithms – Example 2

- Algorithm for Vector Product [Vec-Pro]:
 - How many processors?
- Algorithm [for Vec-Pro]:
 - // Input: Vectors A and B of size n
 - for** all P_i in $i = 1$ to n
 - $C[i] = A[i] * B[i]$
 - // Compute sum $C[i]$ for all i
 - // How do you do this in parallel ?

PRAM Algorithms – Example 3

- Recall the second phase of the algorithm for Vector Product:
 - A list (i.e. an array) of values has to be summed up in one value
- How many parallel tasks are possible?
 - In the first step – *given n values to be added* :
 - $n/2$ additions can be performed in parallel
 - resulting in $n/2$ values to be added
 - which in turn is the same as the original problem

Algorithm Design – Reduction: Example

- SUM (EREW PRAM)

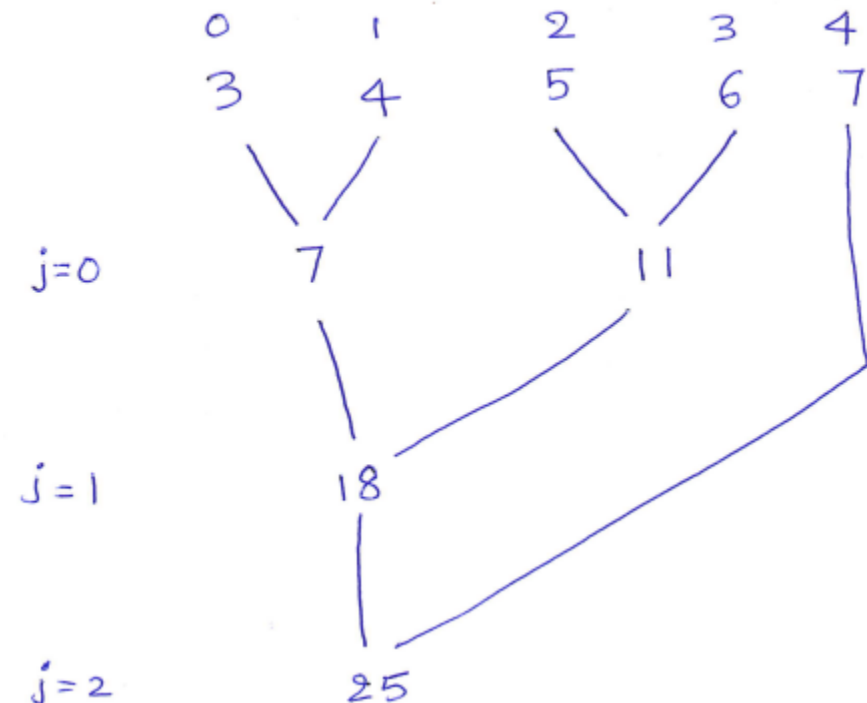
- Precondition: List $A[0..n-1]$ in global memory
- Postcondition: sum $A[0]$ in global memory
- Global variables, A , n , and j

- Parallel reductions

- Parallel summation is an example of divide and conquer as well as parallel reduction
 - At a time two values are added
- A group of n values can be added in $\text{CEIL}(\log n)$ steps

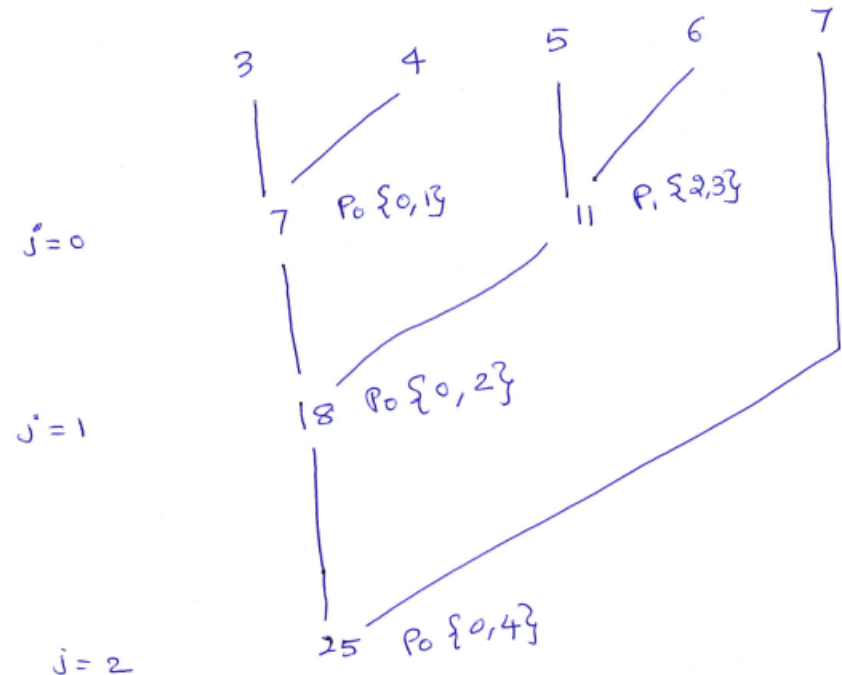
```

for all  $P_i$  where  $0 \leq i \leq \text{FLOOR}(n/2)-1$  {
    for  $j = 0$  to  $\text{CEIL}(\log(n))-1$  {
        ...
    }
}
    
```



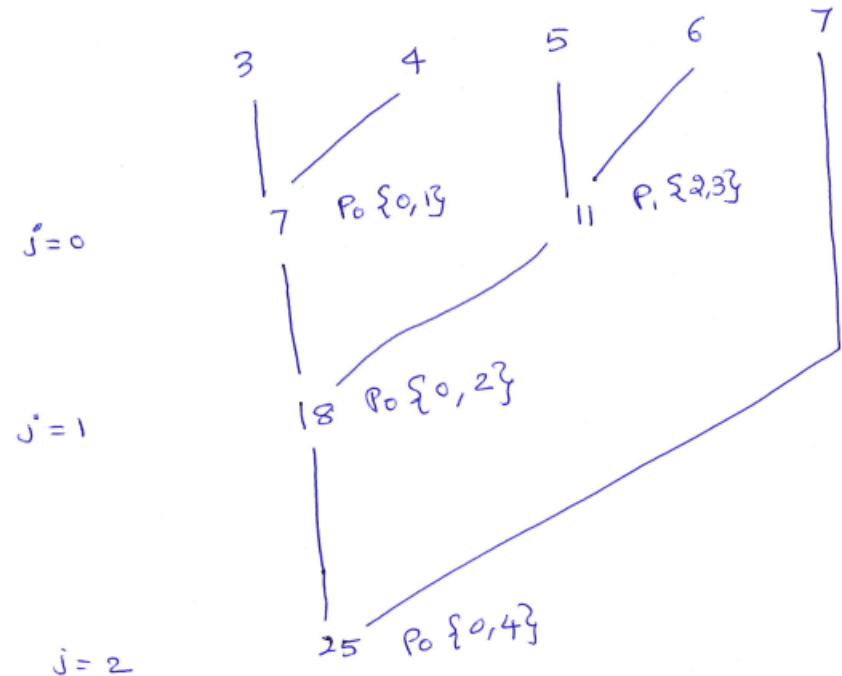
Algorithm Design – Reduction: Example

- How many processes are required?
 - n elements
 - Sum up two values at a time
 - No of processors = $\text{FLOOR}(n)$
- How do we identify elements per process and per step?
 - What should be the relationship between i (rank) and j ?
 - $A[2i] = A[2i] + A[2i + 2^j]$
 - 2^j is the distance between self and (processor holding) other data



Algorithm Design – Reduction: Example

- Are all processors participating in each step?
 - No
 - How do we decide which process should participate?
 - $i \bmod 2^j == 0$
 - $j=0, \{0,0\}$
 - $j=1, \{0,1\}$
 - $j=2, \{0,1\}$



Algorithm Design – Reduction: Example

- Notice that P1 is accessing A[6] which is invalid

- Boundary condition is
 - $2i + 2^j < n$

```

spawn(P0, P1, ... Pk) where k=FLOOR(n/2)-1
for all Pi where 0 <= i <= FLOOR(n/2)-1 {
    for j = 0 to CEIL(log(n))-1 {
        if (i mod 2^j = 0) /* incomplete */ then
            A[2*i] = A[2*i] + A[2*i + 2^j]
    }
}
  
```

```

int A[]={3,4,5,6,7};
//we are using FLOOR(n/2) processes
//NO of processes=5/2=2; i=0,1
//P0,P1
  
```

	P0	P1
j=0	$A[0] = a[0] + A[0 + 2^0]$ $A[0] = A[0] + A[1]$	$A[2] = A[2] + A[2 + 2^0]$ $A[2] = A[2] + A[3]$
j=1	$A[0] = a[0] + A[0 + 2^1]$ $A[0] = A[0] + A[2]$	$A[2] = A[2] + A[2 + 2^1]$ $A[2] = A[2] + A[4]$
j=2	$A[0] = a[0] + A[0 + 2^2]$ $A[0] = A[0] + A[4]$	$A[2] = A[2] + A[2 + 2^2]$ $A[2] = A[2] + A[6]$

Algorithm Design – Reduction: Example

- Are there any concurrent accesses (R/W) in this algorithm?

▪ EREW

```
begin
spawn (P0, P1, ... Pk) where k=FLOOR(n/2)-1
for all Pi where 0 <= i <= FLOOR(n/2)-1 {
  for j = 0 to CEIL(log(n))-1 {
    if (i mod 2^j = 0) and (2*i + 2^j < n) then
      A[2*i] = A[2*i] + A[2*i + 2^j]
  }
}
end
```

```
int A[]={3,4,5,6,7};
```

```
//we are using FLOOR(n/2) processes
```

```
//NO of processes=5/2=2; i=0,1
```

```
//P0,P1
```

	P0
j=0	A[0]=a[0]+A[0+2^0] A[0]=A[0]+A[1]
j=1	A[0]=a[0]+A[0+2^1] A[0]=A[0]+A[2]
j=2	A[0]=a[0]+A[0+2^2] A[0]=A[0]+A[4]

	P1
	A[2]=A[2]+A[2+2^0] A[2]=A[2]+A[3]

Algorithm Design – Reduction: Example

- What is the time complexity?
 - Spawn routine takes $\text{CEIL}(\log(\text{FLOOR}(n/2)))$ steps
 - For loop executes $\text{CEIL}(\log n)$ steps
 - Each iteration has constant time complexity
 - Overall time complexity is $\Theta(\log n)$
 - Given $\text{FLOOR}(n/2)$ processors

```
begin
spawn (P0, P1, ... Pk) where k=FLOOR(n/2)-1
for all Pi where 0 <= i <= FLOOR(n/2)-1 {
    for j = 0 to CEIL(log(n))-1 {
        if (i mod 2^j = 0) and (2*i + 2^j < n) then
            A[2*i] = A[2*i] + A[2*i + 2^j]
    }
}
end
```

Algorithm Design – Speedup and Efficiency

- Speedup:
 - $S = T_{\text{seq}} / T_{\text{par}}$ (generic)
 - $S(p) = T(n,1) / T(n,p)$
 - Ideal speedup should be p .
- Example: Speedup of Summation by Reduction:
 - $S(n/2) = (n/2) / \log(n)$
 - This is less than ideal speedup!
- Efficiency (definition):
 - $E(n,p) = S(p)/p$
 - Ideal efficiency is 1.
- For our example:
 - $E(n,n/2) = (n/2) / \log(n) * (n/2) = 1/\log(n)$

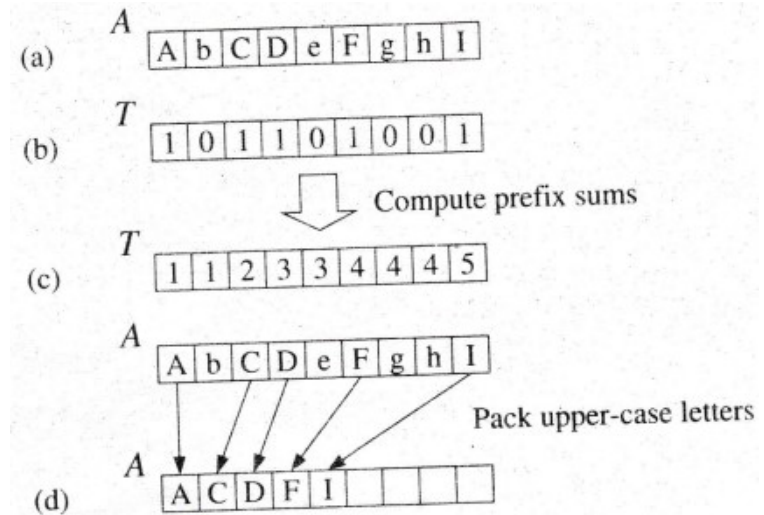
Algorithm Design: Prefix Sum Example

- Given a set of n values $a_1, a_2, a_3, \dots, a_n$, and an associative operator \oplus the prefix sums problem is to compute n quantities:

- a_1
- $a_1 \oplus a_2$
-
- $a_1 \oplus a_2 \oplus a_3 \oplus \dots \oplus a_n$
- Operator

- Prefix sums are also called parallel prefixes or scans

- Have many applications for instance packing elements



- (a) Array A contains both upper and lower-case letters
- (b) Array T contains a 1 for upper, 0 for lower-case
- (c) Array T after prefix-sums
- (d) For each element of A containing upper-case letter, the corresponding element of T is the element's index in the packet array

Algorithm Design: Prefix Sum Example

- PREFIX_SUM (EREW PRAM)
 - Precondition: List $A[0..n-1]$ in global memory
 - Postcondition: Each element $A[i]$ contains its prefix sum
 - Global variables, A , n , and j
- Does it have concurrent accesses?
 - Yes but only reads
- What is the time complexity?
 - $\Theta(\log n)$ given $n - 1$ processors

```
begin
spawn (P0, P1, ... Pk) where k=n-1
for all Pi where 0 <= i <= n-1 {
  for j = 0 to CEIL(log(n))-1 {
    if (i - 2^j >= 0) then
      A[i] = A[i] + A[i - 2^j]
  }
}
end
```

Algorithm Design – Prefix Sum Example

- Speedup:
 - $S = T_{\text{seq}} / T_{\text{par}}$ (generic)
 - $S(p) = T(n,1) / T(n,p)$
 - Ideal speedup should be p .
- Example: Speedup of prefix Summation by Reduction:
 - $S(n/2) = (n-1) / \log(n)$
 - This is less than ideal speedup!
- Efficiency (definition):
 - $E(n,p) = S(p)/p$
 - Ideal efficiency is 1.
- For our prefix sum Example :
 - $E(n,n-1) = (n-1) / \log(n) * (n-1) = 1 / \log(n)$

Parallel Reduction - Template

- Template REDUCE
 - Precondition: Inputs, G , in global memory
 - Postcondition: Result in $G[f(0)]$
 - Global variables: n and j , apart from G
 - begin
 - spawn (P_0, P_1, \dots, P_k) where $k = \text{floor}(n/2) - 1$
 - for all P_i where $0 \leq i \leq \text{floor}(n/2) - 1$ {
 - for $j = 0$ to $\text{ceil}(\log(n)) - 1$ {
 - if $(i \bmod 2^j = 0)$ and $(g(i) < n)$ then
 - $G[f(i)] = G[f(i)]$ BOP $G[g(i)]$
 - }}
- end

BOP is any
binary operation

index of own
data

index of other
data

Parallel Reduction - Template

- Reduction provides a template for
 - parallel execution of
 - *any associative binary operation*
 - extended over a list of values
- Example Instances:
 - Maximum of n values
 - BOP is max
 - Sum of n matrices
 - BOP is matrix addition
 - MergeSort
 - BOP is merging two sorted lists

References

- Chapter 2 from M.J. Quinn, *Parallel Computing : Theory & Practice*, McGraw Hill Inc. 2nd Edition 2002



BITS Pilani
Pilani Campus



Thank You